Simulation Analysis

Sophie Li.cluwartz

2025-05-30

Contents

1 Introduction **1**

2 Research Questions **2**

3 The Simulation **2**
3.1 The data generation parameters 2

4 Data Generating Process **6**
4.1 How the data are generated 6
4.2 How we correlate block-level quantities 8
4.3 How we generate distributions of sizes 8

5 Are any of the estimators ever substantially biased? When? **9**

6 Are any of the estimators generally more precise than others? When? **10**

7 Which estimators produce the best estimates of the treatment effect standard error? **11**

8 Do any of the estimated standard errors target finite versus super population uncertainty? **15**

9 Are problems caused by low degrees of freedom due to estimate instability or bias? **17**

1 Introduction

Blocked cluster randomized controlled trials (CRTs) are a widely used experimental evaluation design in which entire groups (clusters) such as clusters or classrooms are randomly assigned to treatment or control conditions within defined strata or blocks. The effectiveness of the program or treatment is then typically quantified using the average treatment effect (ATE). However, within this relatively straightforward quantity, researchers must make a variety of choices. Researchers need to decide whether to target a person-weighted average or cluster-weighted average; they need to decide whether they are targeting an estimate just for the clusters in the study (finite sample) or an estimate representative of some broader population (super-population); and they need to decide from among different methods or estimators that can be used to estimate the ATE.

This paper contains simulations that complement research examining empirically the implications of different choices researchers make when performing CRTs (CITATION). Simulations, unlike empirical analyses, allow us to compare estimated quantities to their true values, which we cannot observe in empirical contexts. This in turn allows us to evaluate how well different statistical methods perform. Additionally, we can systematically vary data conditions and identify which factors impact method performance.

2 Research Questions

With this in mind, these simulations answer the following research questions:

1. Are any of the estimators ever substantially biased? When?
2. Are any of the estimators generally more precise than others? When?
3. Which estimators produce the best estimates of the treatment effect standard error?
4. Do any of the estimated standard errors target finite versus super-population uncertainty?
5. Are the estimation problems caused by low degrees of freedom due to estimate instability or bias?

3 The Simulation

We ran a series of multifactor simulations, which are simulations where we systematically vary a set of factors to understand how the factors influence the performance of the various estimators, and how the factors interact with each other. In general we are close to full multi-factorial designs, where we vary all factors independently, but we do drop combinations that do not make sense (in particular, we drop scenarios that call for correlations between quantities that do not vary).

3.1 The data generation parameters

The main simulation is a fully factored experiment with the following data generating parameters:

Table 1: Simulation Data Generating Parameters

Parameter	Value(s)
Parameters fixed across all scenarios	
Number of Blocks	5
Average Treatment Effect	.2
Block-Level ICC	.2
Average Number of Observations Per Cluster	20
Number of Block Level Covariates	0
Block-Level (Varies across scenarios)	
Block Type	Same, Messy, Reasonable
Average Number of Clusters per Block	5, 10, 20
Proportion of Clusters Treated	.5, .7
Cluster-Level (Varies across scenarios)	
Correlation of Impact and Size of Cluster	0, .8
Cluster Level Treatment Effect Standard Deviation	0, .3
Cluster Level ICC	0, .2
Variation in cluster size (coefficient of variation)	0, .33, .6
Number of Cluster-Level Covariates	0, 1, 5 [R^2 Always .5]

One of our factors is "block type," which is a set of seven parameters that define how similar or different the blocks are from each other. We did not independently vary these factors so we could control the total number of simulation scenarios. All code is public; others can easily adapt the code to run additional scenarios that are of interest. We have three levels of block type, which we call "Blocks All The Same," "Blocks All Reasonable," and "Blocks All Messy." The "Same" case is where all blocks are identical in their characteristics. The "Reasonable" case is what we might imagine we would see in actual practice. The "Messy" case pushes our simulation to the limits so we can see how various methods fail in extreme circumstance.

Table 2: Block-Type Parameters

Parameter	**Blocks All The Same**	**Blocks All Reasonable**	**Blocks All Messy**
Variation of number of clusters in the blocks (coefficient of variation)	0 (none)	.33 (average)	.85 (big)
Block-Level Average Treatment Effect Standard Deviation	0	.2	.6
Correlation Between Treatment Effect and Cluster Size	0	.2	.8
Correlation Between Average Number of Observations Per Cluster and Total Number of Clusters	0	.3	.8
Correlation of Proportion of Clusters Treated per Block and Impact	0	0	.8
Correlation Between Outcome Intercept and Cluster Size	0	.3	.8

Each scenario would have a specific combination of these parameters. We drop all scenarios with a correlation of a quantity that does not vary (e.g., if there is no treatment effect variation, then we drop scenarios with nonzero correlation between treatment effect and cluster size).

We also ran a separate set of simulations to evaluate how standard errors perform for fully interacted models in matched studies. These simulations were also fully factored and the parameters for those simulations were as follows:

Table 3: Matched Simulation Data Generating Parameters

Parameter	Values
Block-Level Fixed	
Number of Blocks	20
Average Treatment Effect	.2
Block-Level ICC	.2
Average Number of Observations Per Cluster	20
Number of Block Level Covariates	0
Proportion of Clusters Treated	.5
Average Number of Clusters per Block	2
How much do number clusters per block vary (coefficient of variation)?	0
Correlation Between Treatment Effect and Cluster Size	0
Correlation Between Average Number of Observations Per Cluster and Total Number of Clusters	0
Correlation of Proportion of Clusters Treated per Block and Impact	0
Correlation Between Outcome Intercept and Cluster Size	0
Block-Level Varying	
Block-Level Average Treatment Effect Standard Deviation	0, .6, .2
Cluster-Level Fixed	
Number of Cluster-Level Covariates	0
Cluster-Level Varying	
Correlation of Impact and Size of Cluster	0, .8
Cluster Level Treatment Effect Standard Deviation	0, .3
Cluster Level ICC	0, .2
How much does the size of clusters vary (coefficient of variation)?	0, .33, .6

4 Data Generating Process

For most of our simulations, for each scenario defined by sample size, target correlation, predictiveness of covariates, and so forth, we used the following data generating process (DGP):

1. Randomly generate a single set of blocks based on the simulation scenario parameters. Each generated block has a set of characteristics such as the average treatment effect of the block, average cluster size, amount of cluster variation, and so forth.
2. Then, for each simulation iteration:
    a. Generate a sample of clusters within each block according to the block characteristics.
    b. Generate individual-level data within each cluster.
    c. Calculate the true ATEs (both cluster-weighted and individual-weighted) for the generated dataset. These are our target finite-sample estimands.
    d. Assign a proportion of clusters to treatment according to the target block-level proportion.

At this point we have a complete set of generated data. Next, for each simulation iteration we actually conduct a small number of iterations. For each iteration we:

a. Permute treatment across clusters within block to generate a new random assignment.
b. Calculate observed outcomes based on the generated data and the new treatment assignment.
c. Estimate the ATE and SE using all our methods.

The nested permutation approach means that for each fixed set of clusters and blocks, we can calculate finite sample performance. For example, the standard deviation of the ATE estimates within a given dataset across the permutations gives us the true finite sample standard error for that dataset. If we look at the standard deviation of the ATE estimates across all our datasets within a scenario, we have the true cluster-level superpopulation standard error. This standard error *does not* take into account uncertainty at the block level, as the block characteristics are all fixed. So, for example, if there is treatment variation across blocks, our true standard error will not reflect uncertainty due to which blocks were generated, only uncertainty due to which clusters and individuals were generated.

For our target estimands we can use, within each generated dataset, the finite sample estimand. We can also calculate the superpopulation estimand (both person-weighted and cluster-weighted) by averaging the finite sample estimands across all generated datasets within a scenario.

We next outline the data generation parameters that drive our DGP, and then give some further technical detail of how we implemented the above steps.

4.1 How the data are generated

For a given scenario, we start with a set of data generating parameters as described above.

Here is one of the more extreme scenarios from our main simulation:

Table 4: Simulation Scenario Parameters

variable	value
ATE	0.2
ICC.block	0.2
ICC.cluster	0.2
J.bar	10
K	5
R2.cluster	0.5
cv.block	0.85
cv.cluster	0.33
cv.nbar	0.33

variable	value
empirical	TRUE
n.bar	20
nbar.J.cor	0.8
num.cov.cluster	5
num.cov.indiv	0
p.bar	0.7
proptx.impact.cor	0.8
sd.block	0.6
sd.cluster	0.3
size.impact.cor	0.8
size.impact.cor.cluster	0.8
size.intercept.cor	0.8
variable.p	TRUE

Then, to generate the individual blocks, we use a multivariate normal distributions to generate the block-level pairs of average control-side outcome and average treatment effect (in the case of treatment variation). Other block characteristics that are to vary across block (e.g., proportion of clusters treated) are then generated according to various specified distributions. Finally, we induce correlations between block level quantities as described below, if the relevant correlation parameters are nonzero.

When generating block mean outcomes and mean effects, we use the "empirical" setting to make sure our multivariate normal distributions for the amount of cross-block impact variation and so forth equals our set values. This mitigates worry that the simulation results might be dependent on *which* blocks we have generated, as we would expect all sets generated, given a specific set of parameters, would lead to the similar final results. That said, we generate these values to target *block averages*, but if we have variable numbers of clusters per block, the cluster-average effects can deviate from these targets quite substantially, especially if number of clusters and treatment effect are correlated.

Overall, however, the finite nature of the blocks should not be a large concern. In particular, because we have a large number of factors in our experiment, when we look at the main effect of a given factors, we will be averaging across many different scenarios with a different specific set of blocks for each scenario, so any dependencies on the specific set of blocks generated would be ameliorated.

Once generated, we store the block characteristics as a small dataset, which we call a "grid," with each row being a block, and each column a characteristic we need for generating the within-block data (number of clusters, average cluster size, etc.). To illustrate, here is the "grid" of block-level characteristics for the above scenario in our main simulation:

did	tau	alpha	J	p	n.bar	sd.cluster	ICC.cluster	k.clu	R2.cluster	cv.cluster	cor.clu	k.indiv
1	-0.4	-0.5	6	0.7	13	0.3	0.2	5	0.5	0.3	0.8	0
2	-0.4	-0.4	4	0.8	19	0.3	0.2	5	0.5	0.3	0.8	0
3	0.3	0.1	21	0.7	25	0.3	0.2	5	0.5	0.3	0.8	0
4	0.5	0.4	17	0.6	22	0.3	0.2	5	0.5	0.3	0.8	0
5	1.0	0.5	38	0.8	38	0.3	0.2	5	0.5	0.3	0.8	0

and here is a set of descriptive features for data we would get from this fixed set of blocks:

```
## K=5  J.bar=17.20 (13.66, cv=0.79)    n.bar=29.03 (9.29 cv=0.32)  p.bar=0.71 (0.07)
## District stats:
##  block ATE=0.20 (0.60)   block ICC=0.45
##  cor(J,tau)=0.96     cor(ptx,tau)=0.40
##  cor(J,alpha)=0.87       cor(nbar, J)=0.95
```

```
## Cluster stats:
##  school ICC=0.20 (0.00)
##  avg cluster size = 28.9 (range 6--60, sd = 13.0, cv=0.45)
##  cor(n,tau) = 0.80 (across blocks)
## ATE information:
##  Within-block cluster ATE and variation: 0.19 (0.27)
##  Total cluster ATE: 0.56 (0.54)
##  Total individual ATE: 0.75 (0.50)
##
##   * Some stats estimated on 99483 generated observations across 3440 clusters.
```

The block ATE is the simple average of the block-level cluster-average ATEs. It is as we targeted. The bottom line are ATEs for a large dataset generated for this set of blocks: these are our target superpopulation parameters for this scenario, so we can estimate the true cluster- and individidual average ATEs. In this specific example, these are different from each other and from the block average. We also see fairly high correlations of all sorts of features of the data. In this given scenario, the number of clusters per block is varying around 5. The numbers in parentheses are the standard deviations of a given value.

Given our grid of block characteristics we then, for each block in turn, use a multilevel model data generating process (as described more fully in <>) to generate the target number of clusters and individuals.

To keep everything in effect size units, we have a target within-block total variation of $1 - ICC_3$, where ICC_3 is the target intraclass correlation for the blocks. Note that we have $ICC_3 + ICC_2 + \frac{\sigma_\epsilon^2}{R_1^2} = 1$, with ICC_2 being the cluster-level ICC and σ_ϵ^2 the individual-level residual variance.

4.2 How we correlate block-level quantities

Given two vectors of numbers, x and y, that we wish to correlate with target correlation ρ, we do the following:

1. Generate a bivariate standard normal with correlation ρ.
2. Re-order the y values so the pairs of the ranks of x and the ranks of the reordered y correspond to the pairs of ranks of the bivariate distribution.

If x and y are not normal, then this will not produce correlations that have the target correlation on average, but it will generally give high correlations when ρ is high, and negative correlations when ρ is negative. We are correlating based on the normal transform of the ranks, so this is somewhat akin to a rank correlation.

4.3 How we generate distributions of sizes

We have three kinds of sizes to generate: number of clusters per block, number of individuals per cluster, and average size of cluster for each block.

For each of these we use a similar approach. We first specify the desired mean and coefficient of variation (cv). For example, we might want an average of 8 clusters per block with a cv of 0.75. This means the standard deviation of our values would be $sd = cv * mean = 0.75 * 8 = 6$.

We then have three possibilities: either the variance is less than the mean (in which case we can use a binomial distribution), or the variance is larger than the mean (in which case we can use a negative binomial distribution). If the variance is very close to the mean, we use a poisson distribution.

We use this approach for the average cluster size across blocks, even though the average cluster sizes could be fractional.

5 Are any of the estimators ever substantially biased? When?

In 90% of the scenarios we tested, the magnitude of bias was less than 0.01 for all estimators (Figure 1), and the scenarios with bias require the confluence of multiple bias-causing factors (Figure 2). The estimators that theory predicts are unbiased do indeed have nearly no bias (or very low bias) across the full range of simulations. The person-weighted fully interacted estimator had bias as severe as -0.033 in the worst 5% of scenarios tested, which the regression tree shows requires a correlation between the treatment effect and cluster size combined with a small number of clusters per block (Figure 2). For the remaining estimators, except the multilevel fully interacted model, we observe bias as severe as -0.2 in some simulations, primarily driven by the combination of dissimilar blocks and unbalanced treatment assignment of clusters. The multilevel fully interacted model performed well with dissimilar blocks and unbalanced treatment assignment, and instead was most biased when there was a combination of correlation between the number of clusters and the treatment effect, zero cluster-level ICC, and substantial variation in cluster sizes.

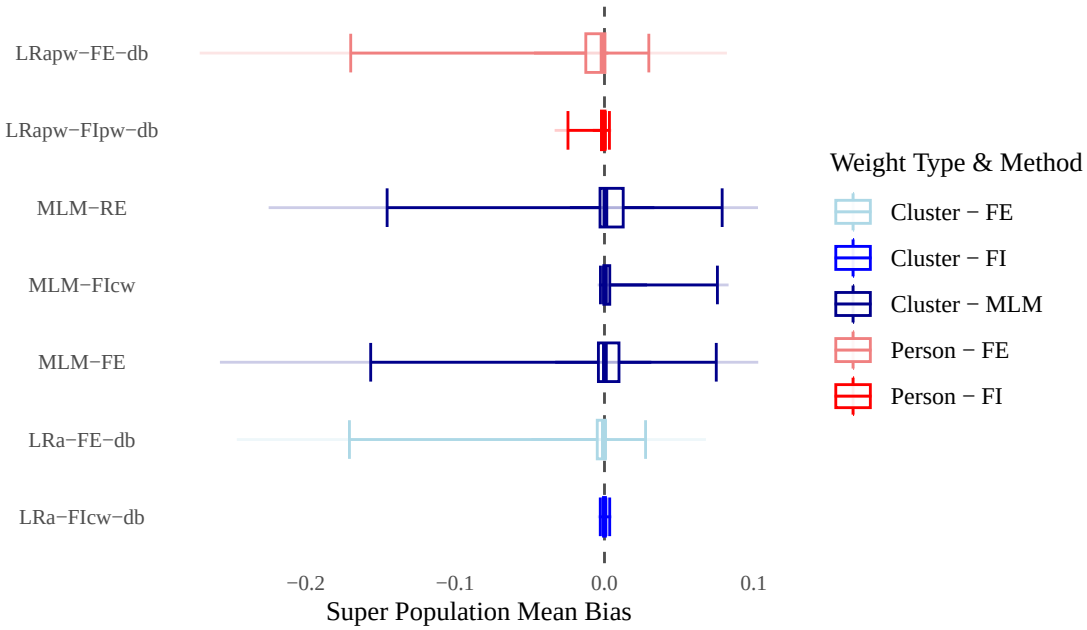


Figure 1: Boxplot of superpopulation mean bias against an estimators targeted estimand across the different scenarios by method. The faint horizontal lines go out to the full range, the bars are the 1% to 99%, and the boxes are the 10% to 90%.

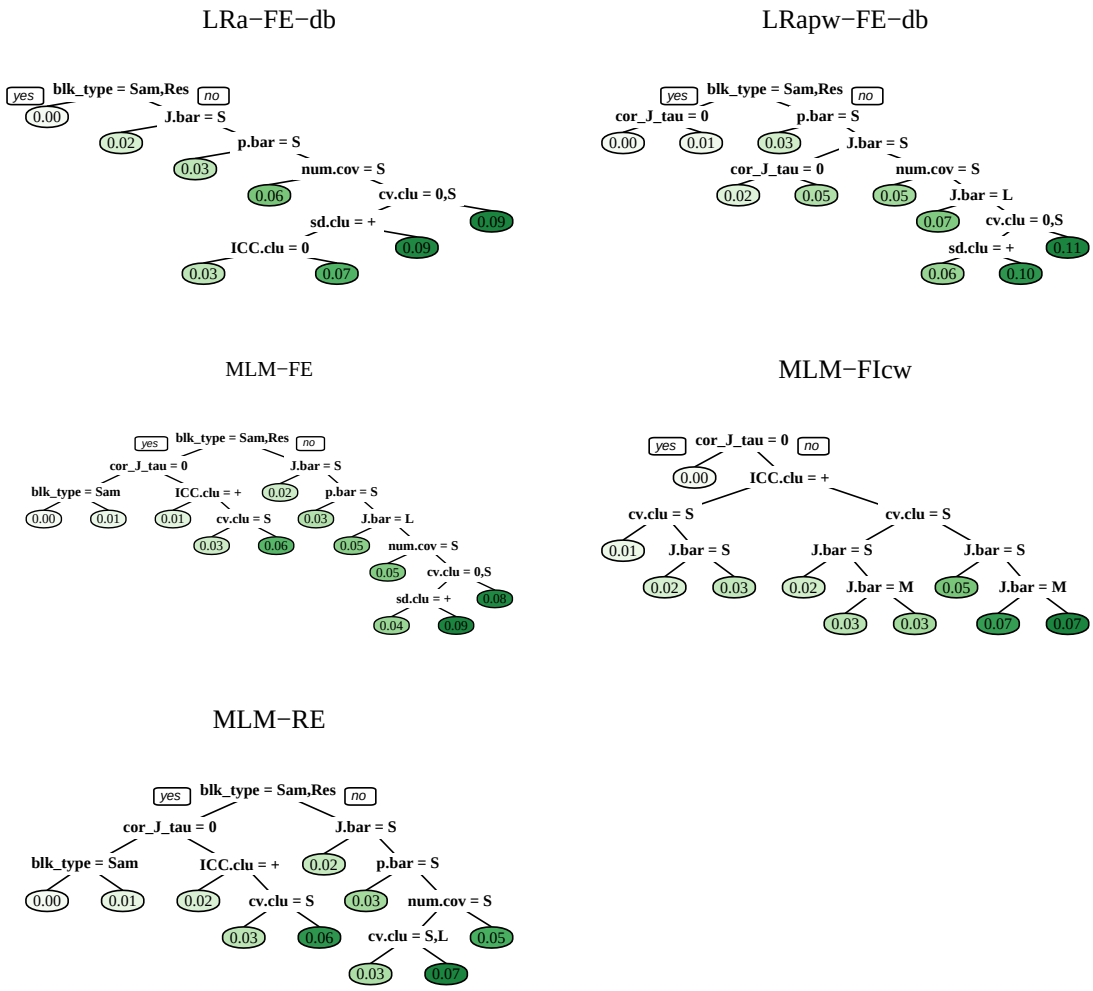


Figure 2: Regression tree by method which predicts the magnitude of the mean bias using different data generating parameters

6 Are any of the estimators generally more precise than others? When?

In general, all the different estimators have comparable precision (Figure 3). Excluding the multilevel estimators, within a given scenario and targeted estimand (person-weighted vs. cluster-weighted), all the estimators have standard errors that are at most 10% larger than the median across methods. The multilevel estimators become more precise than the other estimators when there is a combination of large variation in cluster size and zero cluster ICC; otherwise, their precision is close to that of the other estimators (Figure 4).

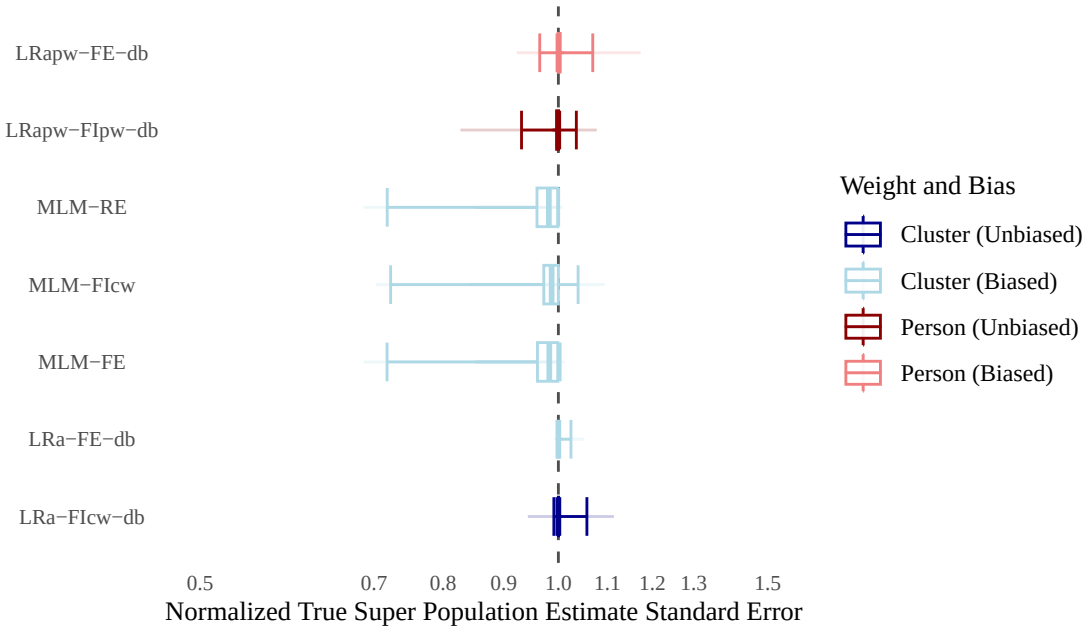


Figure 3: Boxplot of normalized estimate true standard errors across the different scenarios by method. The normalized true standard error is the true standard error for a given method divided by the median true standard error across methods within a given simulation scenario. Therefore each data point is adjusted for

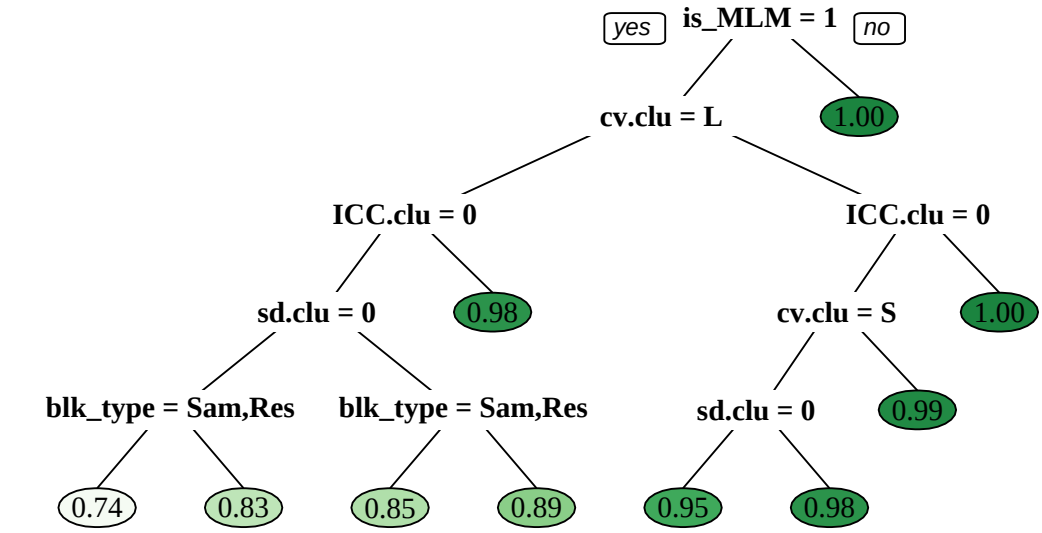


Figure 4: Regression tree by method which predicts the normalized true estimate standard error using different data generating parameters

7 Which estimators produce the best estimates of the treatment effect standard error?

Overall, across the simulations, the standard error estimates for fully interacted estimators, except the fully interacted MLM, are too small (Figure 5). However, across all simulations, the magnitude of the underestimate for the cluster-weighted estimate is very small. For the person-weighted estimators, the standard errors are only substantially underestimated when there is large variation in cluster size combined with either a small number of clusters or large variation across blocks. The only caveat to these results is that almost all the fully interacted estimators are undefined in cases where there are only two clusters in a block. The fully interacted CRVE estimators do have defined standard error estimates in these cases, but these estimates are extremely poor and these estimators should not be used under these circumstances (Figure 6).

On the other hand, the estimated standard errors for the other estimators are on average systematically too high across the simulations. For the MLM, the amount of overestimation is small except when there is a confluence of an unbalanced sample, zero cluster-level ICC, and cross-cluster treatment effect variance (Figure 7). For the fixed effect estimators, the standard error estimates were overestimated by at least 15% in simulations with very dissimilar blocks and no cluster-level ICC (Figure 8).

Consistent with the standard errors being generally well estimated across simulation conditions and estimators, the coverage rate is also close to 95% across simulation conditions and estimators (Figure 9). For the MLM models, sometimes a confluence of multiple simulation conditions (i.e., no variation in cluster size, no cluster-level ICC, and at least three other factors) led the coverage to fall below 90% (Figure 10). Otherwise, there were generally no issues with confidence interval coverage.

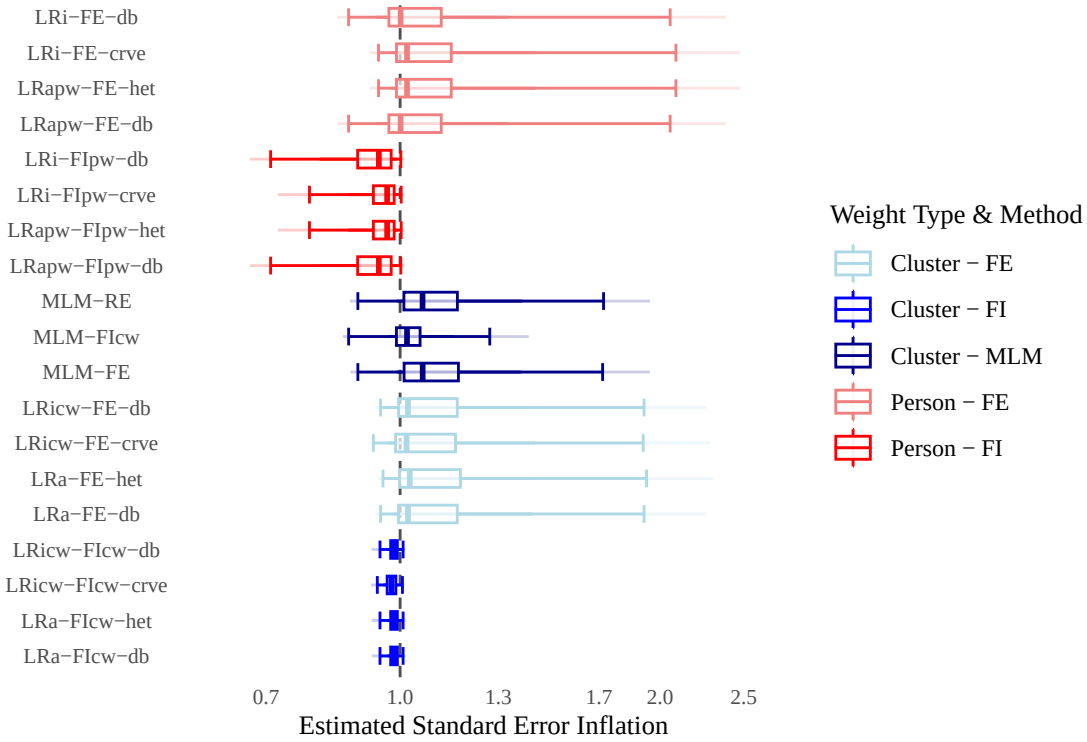


Figure 5: Boxplot of estimated standard error inflation across different simulation scenarios by method. Standard error inflation is measured as the ratio of the estimated standard error over the true super population standard error. Values above 1 indicate overestimation, values below 1 indicate underestimation. The faint horizontal lines go out to the full range, the bars are the 1% to 99%, and the boxes are the 10% to 90%.

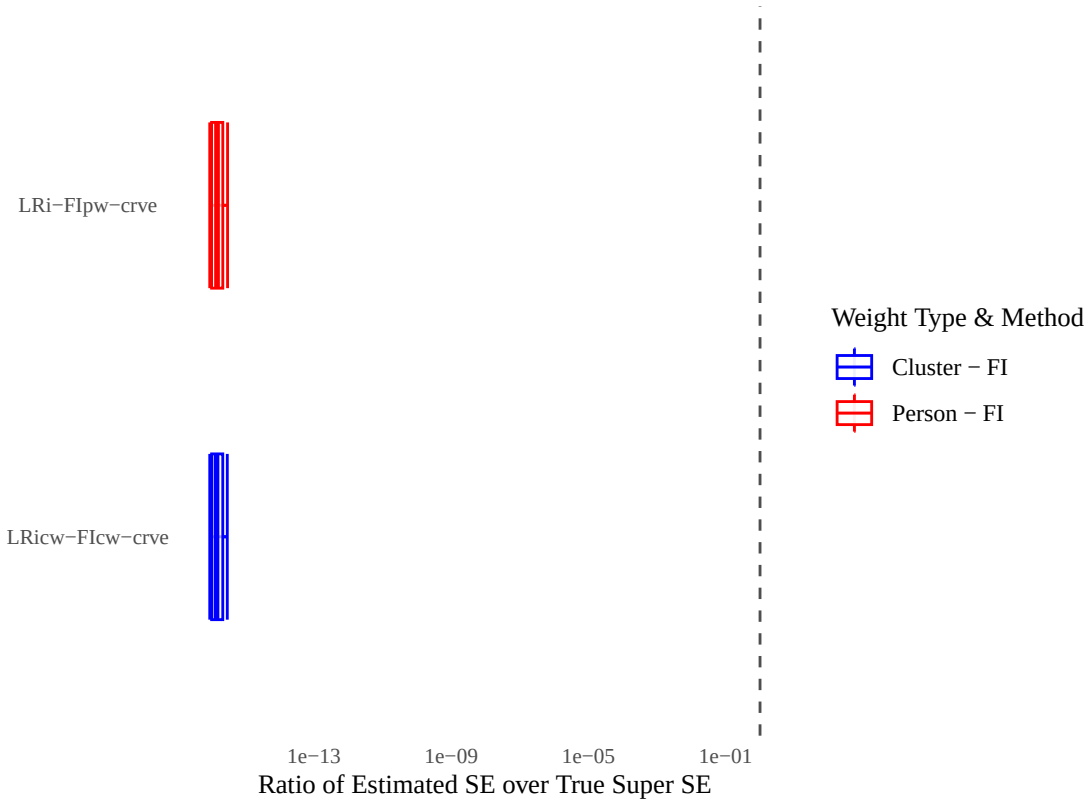


Figure 6: Boxplot of estimated standard error inflation for matched studies for the fully interacted CRVE estimato Standard error inflation is measured as the ratio of the estimated standard error over the true super population standard error. Values above 1 indicate overestimation, values below 1 indicate underestimation. The faint horizon lines go out to the full range, the bars are the 1% to 99%, and the boxes are the 10% to 90%.

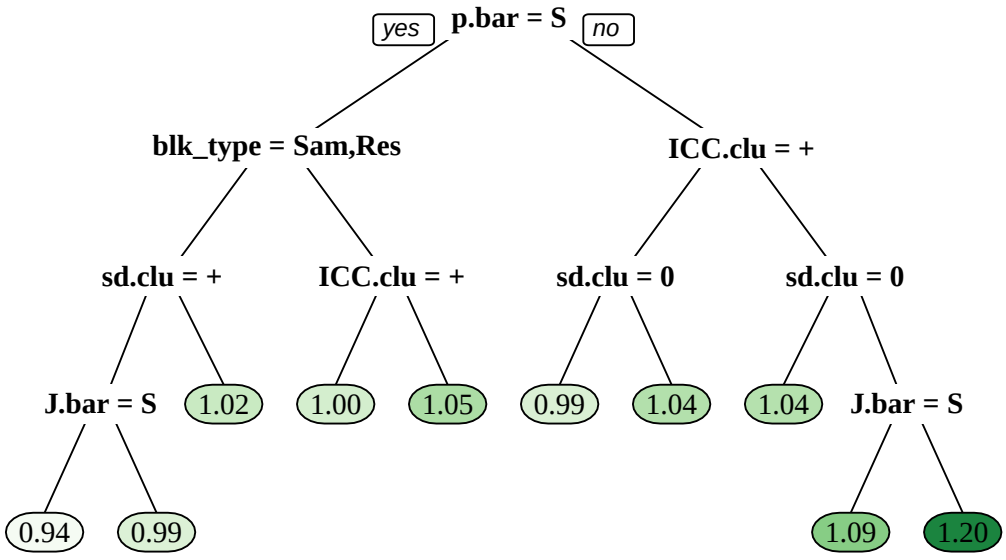


Figure 7: Regression tree which predicts the standard error inflation using different data generating parameters for the multi−level model methods

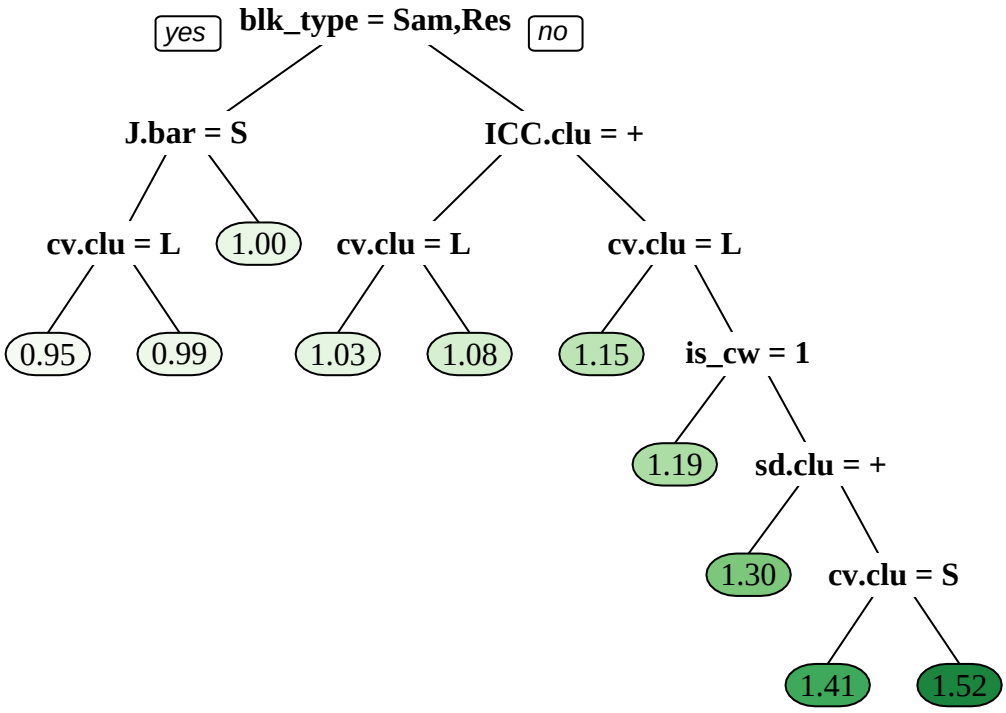


Figure 8: Regression tree which predicts the standard error inflation using different data generating parameters for the fixed effect methods

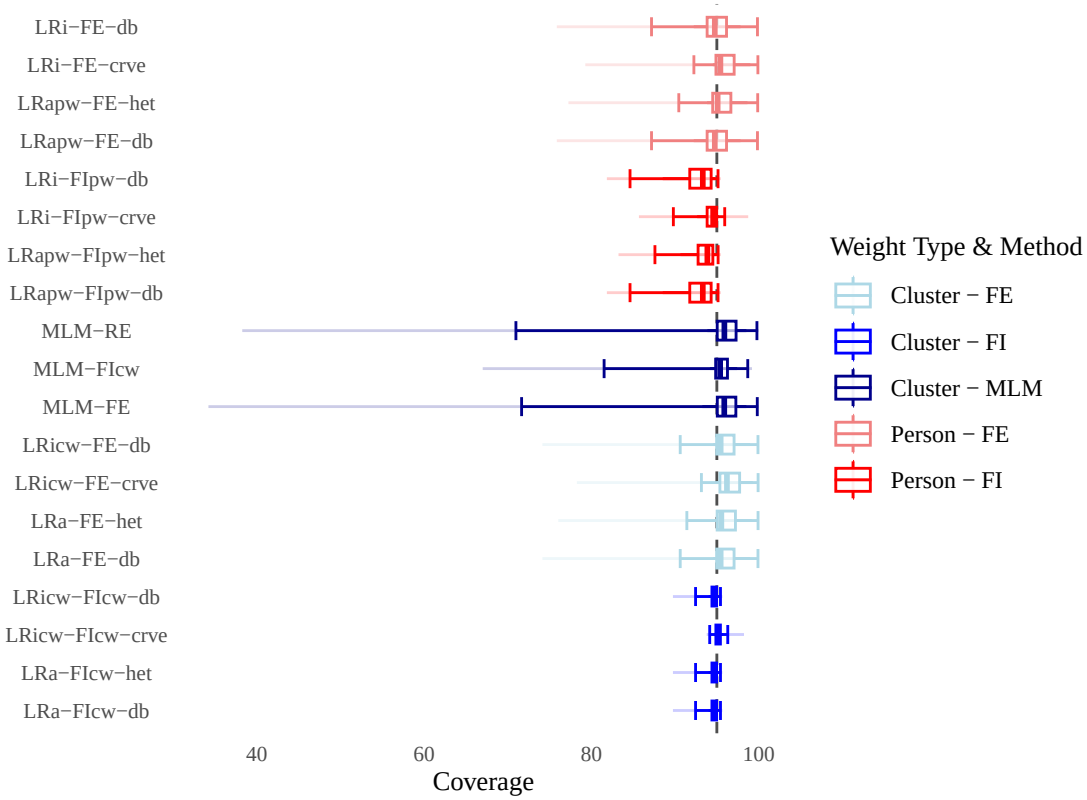


Figure 9: Boxplot of coverage rates across different scenarios by method. The horizontal line indicates the target 95% coverage rate. The faint horizontal lines go out to the full range, the bars are the 1% to 99%, and the boxes are the 10% to 90%.

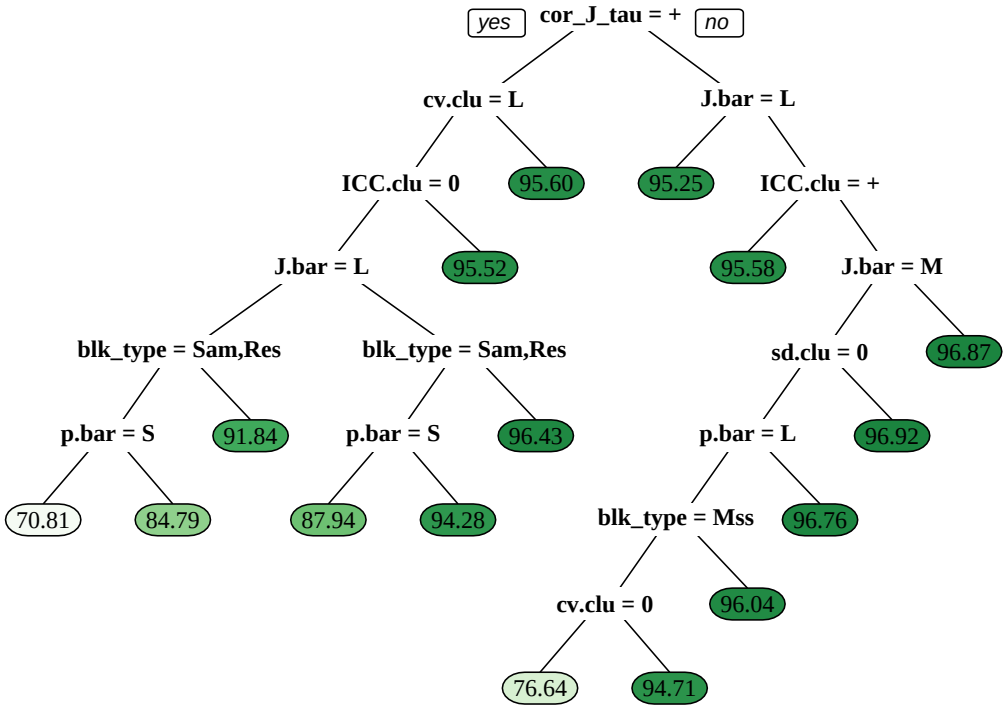


Figure 10: Regression tree which predicts the coverage rates using different data generating parameters for multi−level model methods

8 Do any of the estimated standard errors target finite versus super population uncertainty?

The simulations demonstrate that the standard errors are targeting the super-population standard errors. The true super-population standard errors are only noticeably different from the true finite population standard errors when the cluster-level ICC is zero and there is cross-cluster treatment effect variance (Figure 11). Looking at the subset of simulations where that is true shows that the standard error inflation is consistently closer to one for the true super-population standard errors than for the true finite population standard errors (Figure 12).

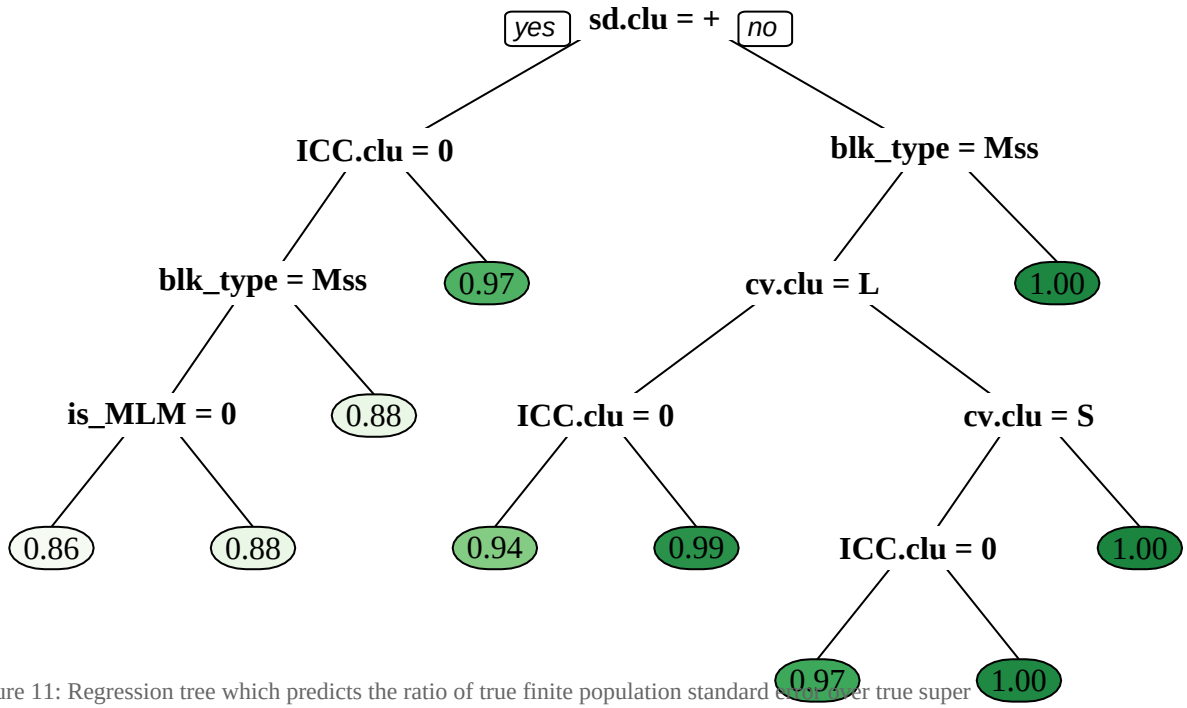


Figure 11: Regression tree which predicts the ratio of true finite population standard error over true super population standard error using different data generating parameters. These two quantities are the same when this ratio is one.

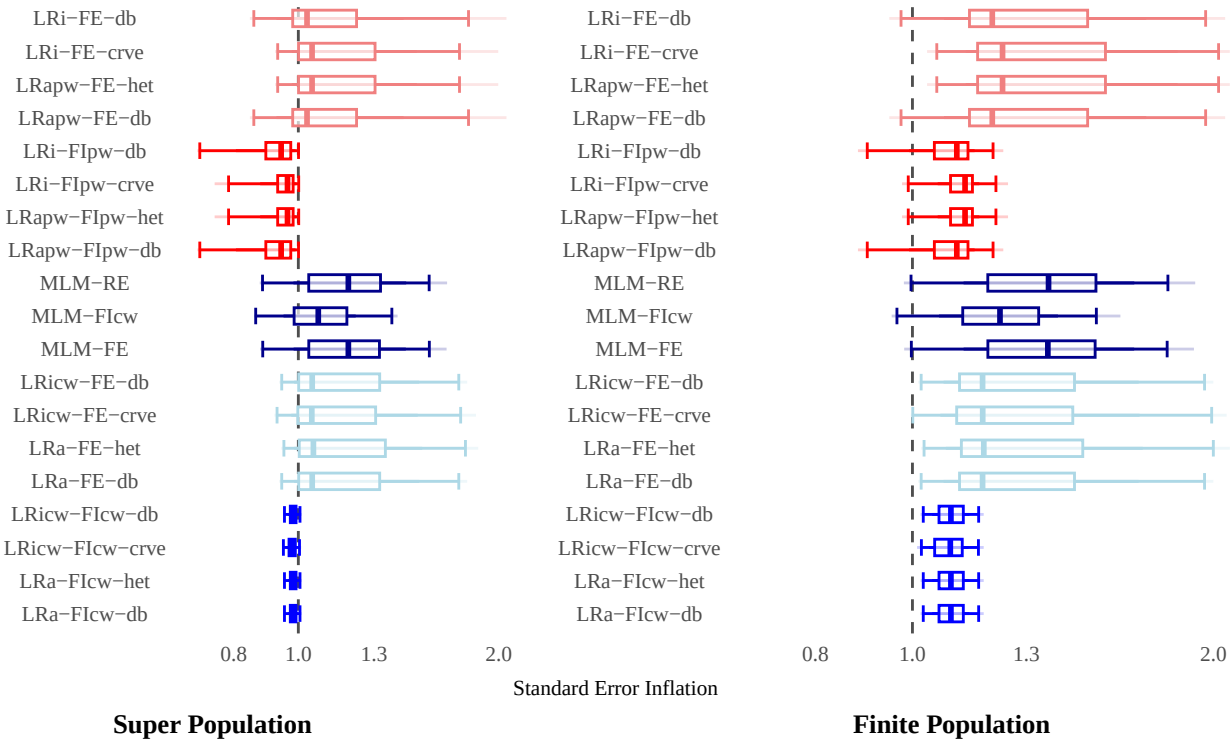


Figure 12: Boxplot of estimated standard error inflation for across simulation scenarios with no cluster−level ICC and non−zero treatment effect heterogeneity by method, comparing againt the true super population and true finite population standard errors. Standard error inflation is measured as the ratio of the estimated standard error over the true standard error. The faint horizontal lines go out to the full range, the bars are the 1% to 99%, and the boxes are the 10% to 90%.

9 Are problems caused by low degrees of freedom due to estimate instability or bias?

Simulations show that the problems with estimates in studies with low degrees of freedom are driven by estimate instability and not estimate bias. We see that as the degrees of freedom get smaller, the estimate instability gets larger (Figure 13). On the other hand, we see no relationship between mean bias and degrees of freedom (Figure 14).

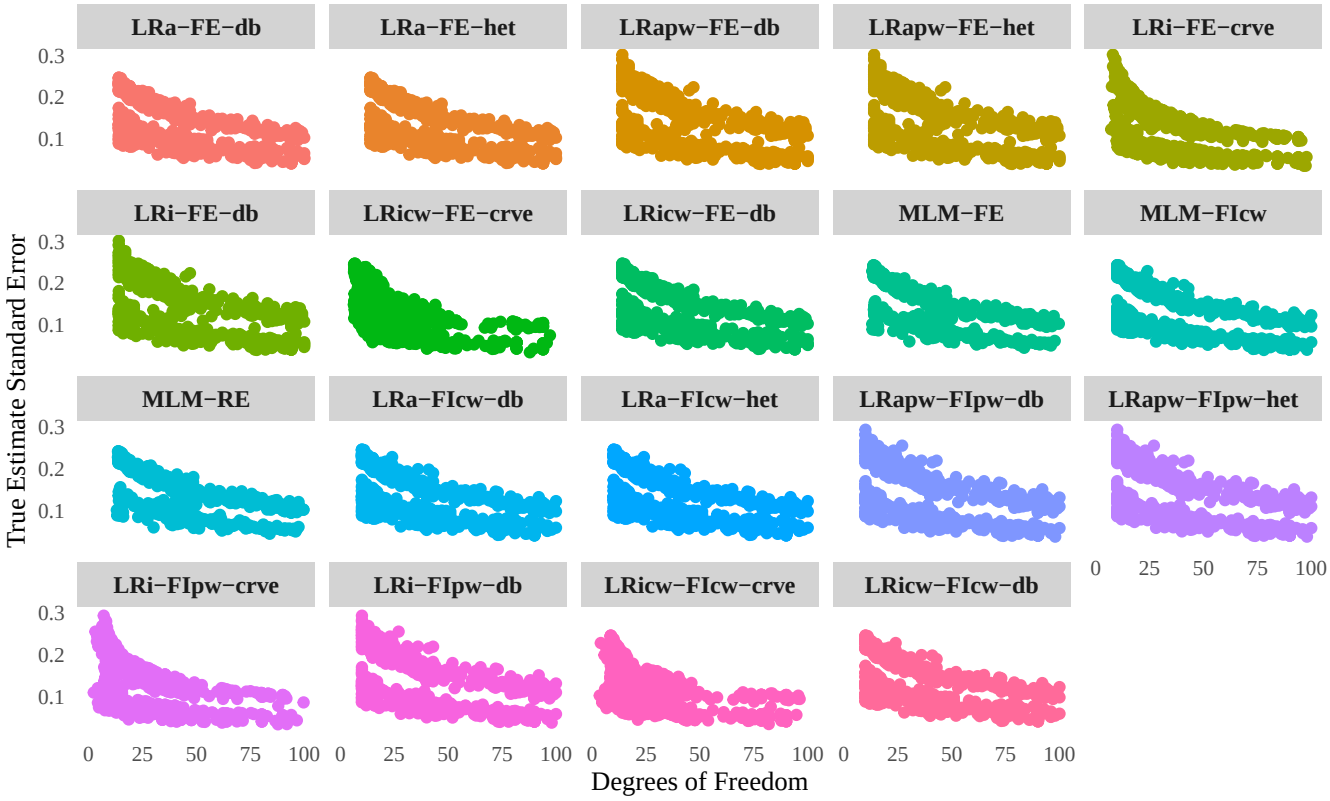


Figure 13: Scatter plot of true estimate standard error versus degrees of freedom by method. Each panel shows the relationship for a different estimation method.

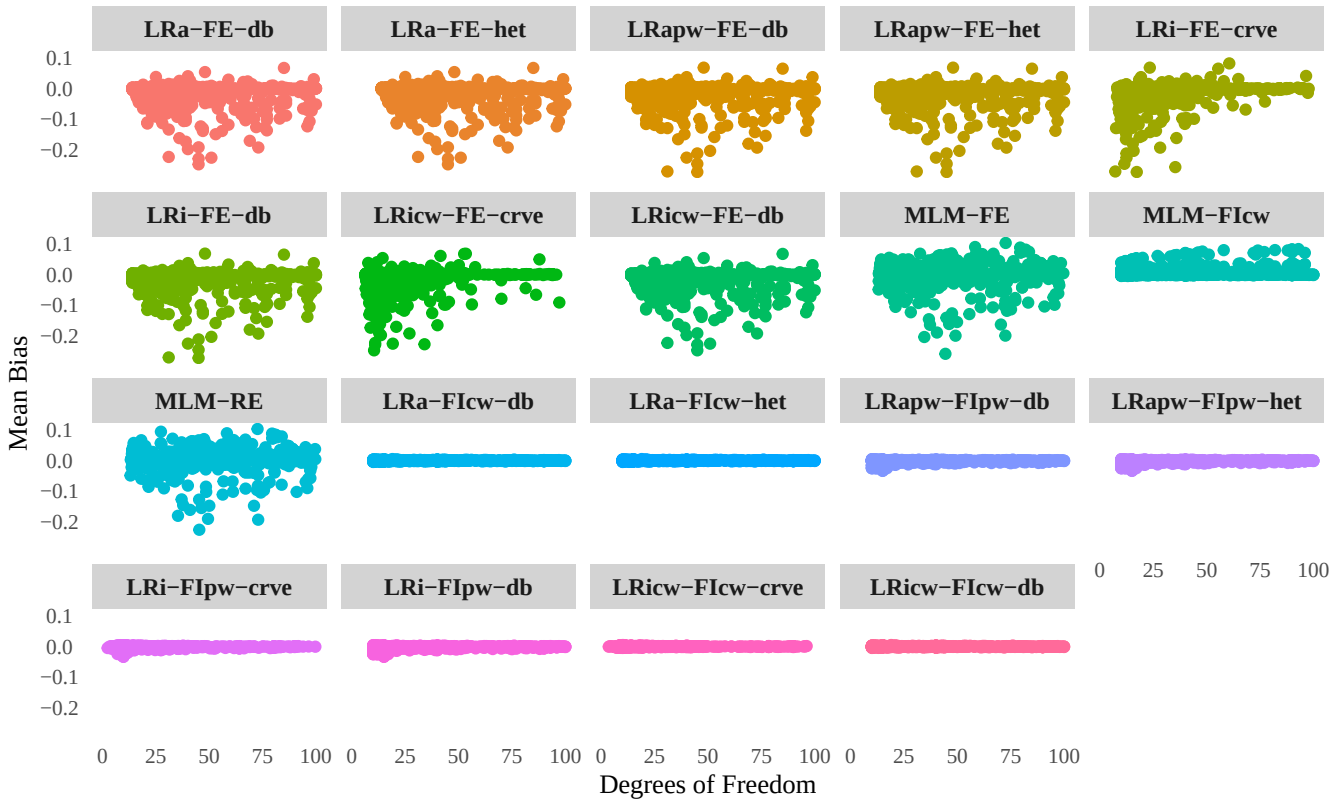


Figure 14: Scatter plot of mean bias versus degrees of freedom by method. Each panel shows the relationship for a different estimation method.